

LLM API 全量速度测试报告

测试时间: 2026-05-31 09:53:27 | 轮次: 3轮/模型 | max_tokens: 800 | temperature: 0.0

提供商: NVIDIA NIM (key1+key2) | v36 中转站 | 讯型AI 中转站 | DeepSeek 直连 | GLM 直连

🏆 速度排名 (共 27 个模型)

排名	模型	提供商	平均耗时	R1	R2	R3	波动
1	⚡️⚡️ deepseek-chat (直连)	直连	1.3s	1.2s	1.5s	1.2s	🟢
2	⚡️⚡️ gpt-5.4-nano (讯型AI)	讯型AI	1.7s	1.5s	2.1s	1.6s	🟡
3	⚡️ deepseek-chat (讯型AI)	讯型AI	2.1s	1.8s	2.7s	1.8s	🟡
4	⚡️ kimi-k2.6 (讯型AI)	讯型AI	2.1s	2.2s	2.4s	1.9s	🟢
5	⚡️ gpt-4o-mini (v36)	v36	2.2s	2.0s	2.8s	1.8s	🟡
6	⚡️ claude-sonnet-4-6 (讯型AI)	讯型AI	2.3s	2.1s	2.6s	2.3s	🟡
7	⚡️ claude-opus-4-7 (讯型AI)	讯型AI	2.3s	2.2s	2.3s	2.5s	🟢
8	⚡️ gpt-4.1-nano (v36)	v36	2.4s	2.2s	2.5s	2.3s	🟢
9	⚡️ gpt-5.2 (v36)	v36	2.4s	2.2s	2.6s	2.4s	🟢
10	⚡️ claude-haiku-4-5 (讯型AI)	讯型AI	2.5s	2.2s	2.7s	2.5s	🟢
11	⚡️ gpt-5.4 (v36)	v36	2.8s	2.9s	2.6s	2.8s	🟢
12	⚡️ DS V4 Pro (NIM2)	NIM2	3.0s	4.1s	3.4s	1.4s	🟠
13	✅ gpt-5.1 (v36)	v36	3.0s	3.3s	2.9s	2.9s	🟢
14	✅ Mistral Large (NIM2)	NIM2	3.5s	1.5s	2.2s	6.8s	🔴
15	✅ deepseek-v4-flash (讯型AI)	讯型AI	3.8s	3.7s	3.1s	4.7s	🟡
16	✅ o3-mini (v36)	v36	4.2s	3.9s	3.7s	4.9s	🟡
17	✅ gemini-3.5-flash (讯型AI)	讯型AI	4.2s	3.0s	3.3s	6.2s	🟠
18	✅ gpt-5.5 (讯型AI)	讯型AI	4.3s	4.5s	4.5s	3.7s	🟡
19	✅ gpt-5-mini (v36)	v36	4.6s	4.6s	4.4s	4.6s	🟢
20	✅ o4-mini (v36)	v36	4.8s	4.9s	3.5s	6.0s	🟠
21	✅ Mistral Large (NIM1)	NIM1	5.2s	1.7s	11.6s	2.2s	🔴
22	✅ MiniMax-M2.7 (讯型AI)	讯型AI	5.5s	5.0s	7.5s	3.9s	🟠
23	✅ DS V4 Flash (NIM2)	NIM2	6.3s	1.1s	11.6s	6.2s	🔴
24	✅ glm-4-flash (直连)	直连	6.3s	1.7s	3.7s	13.6s	🔴
25	✅ glm-5 (讯型AI)	讯型AI	8.0s	7.4s	8.1s	8.4s	🟡
26	🟡 MiniMax M2.7 (NIM1)	NIM1	9.1s	7.5s	6.8s	13.0s	🔴
27	🟡 Qwen 3.5 (NIM1)	NIM1	10.1s	15.2s	5.0s	-	🔴

讯型AI (11 个模型, 平均 3.5s)

- gpt-5.4-nano (讯型AI): 1.7s (成功率 100%)
- deepseek-chat (讯型AI): 2.1s (成功率 100%)
- kimi-k2.6 (讯型AI): 2.1s (成功率 100%)
- claude-sonnet-4-6 (讯型AI): 2.3s (成功率 100%)
- claude-opus-4-7 (讯型AI): 2.3s (成功率 100%)
- claude-haiku-4-5 (讯型AI): 2.5s (成功率 100%)
- deepseek-v4-flash (讯型AI): 3.8s (成功率 100%)
- gemini-3.5-flash (讯型AI): 4.2s (成功率 100%)
- gpt-5.5 (讯型AI): 4.3s (成功率 100%)
- MiniMax-M2.7 (讯型AI): 5.5s (成功率 100%)
- glm-5 (讯型AI): 8.0s (成功率 100%)

v36 (8 个模型, 平均 3.3s)

- gpt-4o-mini (v36): 2.2s (成功率 100%)
- gpt-4.1-nano (v36): 2.4s (成功率 100%)
- gpt-5.2 (v36): 2.4s (成功率 100%)
- gpt-5.4 (v36): 2.8s (成功率 100%)
- gpt-5.1 (v36): 3.0s (成功率 100%)
- o3-mini (v36): 4.2s (成功率 100%)
- gpt-5-mini (v36): 4.6s (成功率 100%)
- o4-mini (v36): 4.8s (成功率 100%)

NIM1 (3 个模型, 平均 8.1s)

- Mistral Large (NIM1): 5.2s (成功率 100%)
- MiniMax M2.7 (NIM1): 9.1s (成功率 100%)
- Qwen 3.5 (NIM1): 10.1s (成功率 67%)

NIM2 (3 个模型, 平均 4.3s)

- DS V4 Pro (NIM2): 3.0s (成功率 100%)
- Mistral Large (NIM2): 3.5s (成功率 100%)
- DS V4 Flash (NIM2): 6.3s (成功率 100%)

直连 (2 个模型, 平均 3.8s)

- deepseek-chat (直连): 1.3s (成功率 100%)
- glm-4-flash (直连): 6.3s (成功率 100%)

实验推荐

第一梯队 — 极快稳定 (<2s)

- deepseek-chat (直连) (直连) — 1.3s, 波动0.3s
- gpt-5.4-nano (讯型AI) (讯型AI) — 1.7s, 波动0.6s

第二梯队 — 快速稳定 (2-3s)

- **deepseek-chat (讯型AI)** (讯型AI) — 2.1s
- **kimi-k2.6 (讯型AI)** (讯型AI) — 2.1s
- **gpt-4o-mini (v36)** (v36) — 2.2s
- **claude-sonnet-4-6 (讯型AI)** (讯型AI) — 2.3s
- **claude-opus-4-7 (讯型AI)** (讯型AI) — 2.3s
- **gpt-4.1-nano (v36)** (v36) — 2.4s
- **gpt-5.2 (v36)** (v36) — 2.4s
- **claude-haiku-4-5 (讯型AI)** (讯型AI) — 2.5s
- **gpt-5.4 (v36)** (v36) — 2.8s
- **DS V4 Pro (NIM2)** (NIM2) — 3.0s

第三梯队 — 可用 (3-8s)

- **gpt-5.1 (v36)** (v36) — 3.0s
- **Mistral Large (NIM2)** (NIM2) — 3.5s
- **deepseek-v4-flash (讯型AI)** (讯型AI) — 3.8s
- **o3-mini (v36)** (v36) — 4.2s
- **gemini-3.5-flash (讯型AI)** (讯型AI) — 4.2s
- **gpt-5.5 (讯型AI)** (讯型AI) — 4.3s
- **gpt-5-mini (v36)** (v36) — 4.6s
- **o4-mini (v36)** (v36) — 4.8s
- **Mistral Large (NIM1)** (NIM1) — 5.2s
- **MiniMax-M2.7 (讯型AI)** (讯型AI) — 5.5s
- **DS V4 Flash (NIM2)** (NIM2) — 6.3s
- **glm-4-flash (直连)** (直连) — 6.3s
- **glm-5 (讯型AI)** (讯型AI) — 8.0s

不推荐 (>8s 或 不稳定)

- **MiniMax M2.7 (NIM1)** (NIM1) — 9.1s
- **Qwen 3.5 (NIM1)** (NIM1) — 10.1s